

# Deep learning and representation


---

## AI ENGINEERING: GENAI AND AGENTIC SYSTEMS

Today you build Project 2, `nw/semantic`: a transformer encoder fine-tuned with LoRA for tags and priority, exported to ONNX with parity proof and an int8 graph, benchmarked against Project 1, and put through the same loop: a gate that reads three files and chooses the graph to serve, a six-step pipeline, a registry version you approve, serving from the registry, and a retraining run with a tightened bar.

# Project 1 orders the queue and cannot say what a ticket is about

---



Northwind wants tags, several per ticket, and similar past tickets in front of the agent. Project 2 is DistilBERT with LoRA adapters, exported to ONNX, benchmarked against Project 1 on the same test split, and gated by the same loop.

# Today you fine-tune a small encoder and hold it to Project 1's bar

---

- 1 Device, memory, the data contract
- 2 Read the loop and see the tests fail
- 3 Freeze the encoder, inject LoRA
- 4 The loss and the training step
- 5 Per-tag thresholds, then train
- 6 Track and compare runs
- 7 Checkpoints that resume
- 8 Export, quantise, prove parity
- 9 The benchmark table
- 10 The gate moves latest; the card
- 11 The semantic pipeline
- 12 Approve and serve from the registry
- 13 Service, third drift signal
- 14 Backtest the next version
- 15 Retrain with a tightened bar
- 16 German error analysis

The deliverable is a benchmark table and a gate decision on it, not a personal best on F1.

# Attention replaced recurrence in 2017, and today's models are built on it

---

## Neural networks

Layers learn their own features from raw input. Project 1's TF-IDF features were designed by hand.

## 2017, the transformer

Attention alone, no recurrence: a new best on translation after 3.5 days on eight GPUs.

## 2018, BERT

An encoder: pretrain on unlabelled text, then fine-tune with one extra output layer.

## Decoders, the LLMs

The same block trained to predict the next token. Generation, and Project 3.

Project 2 is the third card: a pretrained encoder, fine-tuned for Northwind's tags and priority. It reads a ticket; it does not write one.

# Prompt first; fine-tune when the task is narrow, frequent and yours

| Approach    | What changes                                    | Who pays for it | Here            |
|-------------|-------------------------------------------------|-----------------|-----------------|
| Pretraining | every weight, from vast unlabelled text         | a lab, once     | you download it |
| Fine-tuning | weights, for one task, from your labels         | you, per task   | Project 2       |
| Prompting   | nothing in the model: instructions and examples | you, per call   | Project 3       |

**Fine-tune when** you have more examples than fit in a prompt, need lower cost and latency at scale, hold private training data, or want a smaller model to do one task well. Otherwise, with evals in place, a prompt may be all you need.

# LoRA trains a sliver of the weights, and learns less but forgets less

---

# 10,000x

fewer trainable parameters than full fine-tuning of a 175B model, with 3x less GPU memory

- Freeze the pretrained weights; train small low-rank matrices beside them: adapters
- QLoRA fine-tunes a 65B model on one 48 GB GPU with 4-bit base weights
- The trade: LoRA learns the new task less well than full fine-tuning, and keeps more of what the model already knew
- Project 2 puts adapters on the encoder's attention and trains the two heads

# Distillation and quantization make a model cheap enough to serve

---

## Distillation: a small student learns from a big teacher

- DistilBERT: 40 percent smaller, 60 percent faster, 97 percent of BERT's language understanding. Project 2's base
- DeepSeek distilled its reasoning model into open models from 1.5B to 70B; the 32B one beat o1-mini on AIME 2024

## Quantization: fewer bits per weight

- int8 halves inference memory with no loss, up to 175B
- 4-bit methods (GPTQ, AWQ) run over 3x faster than FP16
- ONNX Runtime recommends dynamic int8 for transformers: Project 2's export

## "It does not fit" is usually the optimiser, not the model

| Held per parameter, fp32 with Adam | Bytes     |
|------------------------------------|-----------|
| Weights                            | 4         |
| Gradients                          | 4         |
| Two optimiser moments              | 8         |
| <b>Full fine-tuning</b>            | <b>16</b> |

# about 1 GB

66 million DistilBERT parameters at 16 bytes each, before a single activation

Activations scale with batch size times sequence length times hidden size times layers. Write the number down before you touch a model.

# Freeze the encoder and train under one percent of it

---

```
for p in model.encoder.parameters():  
    p.requires_grad = False  
cfg = LoraConfig(r=8, lora_alpha=16,  
                lora_dropout=0.05,  
                target_modules=["q_lin", "v_lin"],  
                bias="none")  
model.encoder = get_peft_model(model.encoder, cfg)
```

## 0.29%

192,058 of 66,554,938 parameters train: the adapters on query and value, and the two heads

# Multi-label is not multi-class

|                          | Tags                                        | Priority                                   |
|--------------------------|---------------------------------------------|--------------------------------------------|
| Can two be true at once? | Yes: "SSO" and "Outage"                     | No                                         |
| Loss                     | binary cross - entropy per tag              | cross - entropy with class weights         |
| Decision                 | one threshold per tag, chosen on validation | argmax                                     |
| Wrong choice looks like  | softmax across tags: exactly one fires      | a flat 0.5 under - predicts every rare tag |

Priority here has no threshold rule: this head does not protect P0, and the benchmark will show it.

# Training ends with a candidate, not a decision

# 0.742

P0 recall on the `test:` line, 23 of 31. 0.839 on the last validation line. Same model, same argmax

- 2,000 ticket subset, six epochs, about 100 seconds on Apple silicon
- Test: tag micro-F1 0.541, tag macro-F1 0.609, priority macro-F1 0.693
- `candidate <version> registered, not promoted` : this gate reads parity and the int8 row, which do not exist yet
- `make runs-semantic` : one row, the same numbers, the seconds and the device

# ONNX is the serving boundary

---

- 1 Named inputs and outputs, dynamic axes on batch and sequence, opset 17, `dynamo=False`
- 2 Parity on 20 held-out tickets: fp32 max logit difference  $1.0e-5$ ; the gate holds it to  $1e-4$
- 3 Dynamic int8: 265.7 MB to 66.8 MB, ratio 3.98, judged by task metrics, not logits

## Where ONNX leaks

Custom operators, data-dependent control flow, anything the tracer turns into a constant.

If parity fails only on the second ticket of a batch, the axes are fixed.

# Project 2 adds tags at eighteen times Project 1's latency, and misses more emergencies

---

| Model                                     | Priority macro - F1 | P0 recall | Tag micro - F1 | p50 ms | p95 ms | Size MB |
|-------------------------------------------|---------------------|-----------|----------------|--------|--------|---------|
| Project 1: TF-IDF and logistic regression | 0.681               | 0.903     | n/a            | 1.7    | 2.0    | 6.9     |
| Project 2: PyTorch fp32                   | 0.693               | 0.742     | 0.541          | 20.2   | 27.2   | 268.0   |
| Project 2: ONNX fp32                      | 0.693               | 0.742     | 0.541          | 30.7   | 45.6   | 265.7   |
| Project 2: ONNX int8                      | 0.666               | 0.645     | 0.524          | 11.6   | 16.9   | 66.8    |

Six epochs, Apple silicon, test split, batch size one, measured 2026-09-30. Your latencies are your machine's; the quality columns are not.

# The gate reads three files and chooses the graph it will serve

| Bar                            | Value                                        | fp32         | int8                    |
|--------------------------------|----------------------------------------------|--------------|-------------------------|
| Tag micro-F1, served graph     | at least 0.50                                | 0.541        | 0.524                   |
| P0 recall, served graph        | at least 0.70                                | 0.742, 23/31 | <b>0.645, 20/31</b>     |
| fp32 ONNX parity               | at most 1e-4                                 | 1.0e-5       |                         |
| int8 against fp32              | drops at most 0.03, at most 1 more missed P0 |              | 0.027, 0.017, <b>+3</b> |
| int8 p95 latency               | at most 100 ms                               |              | 16.9                    |
| Against production, both heads | drop at most 0.02                            | 0.000, 0.000 |                         |

**PROMOTED**, serves the fp32 graph, written to `serving.json`, which the service reads.

# The pipeline puts three steps between training and the gate

| Step      | Adds                                                                            |
|-----------|---------------------------------------------------------------------------------|
| data prep | Project 1's contract, then the tag checks                                       |
| train     | the candidate, as <code>make train-semantic</code> makes it                     |
| export    | the graphs, int8, parity                                                        |
| benchmark | <code>benchmark.json</code> , training a Project 1 reference if none is at hand |
| gate      | the decision, <code>passed_int</code> for SageMaker's ConditionStep             |
| register  | the candidate, only if the gate passed                                          |

## Same loop, one more pipeline

The same Kubeflow YAML on the local runner and on Pipelines; a SageMaker definition over the same step code on AWS; Azure Machine Learning pipelines on Azure.

The cloud instances have no accelerator: the run takes many times longer than on a laptop. Submit and go on.

# Approve what will be served, and say why in the reason

The registry version carries the fp32 and int8 metrics side by side, and the served format.

Today that is fp32: `/version` reports `format: model.onnx`.

The reason names the tag score, the int8 refusal and the German line.

## AWS

The approval event creates `northwind-<handle>-semantic` on Serverless Inference; solo mode deploys directly.

## Google Cloud, Azure, Local

A new Cloud Run revision reading `NW_MODEL_URI`; your semantic managed online endpoint; `models:/northwind-solo-semantic/1` on :8002.

# The tag rate says the model is firing on noise

---

# 3.85

tag rate PSI after 200 forty-character tickets. Text length 8.87, as in Project 1; priority 4.91; level alert

- Text length and priority reuse Project 1's monitor unchanged, with the same 200-request minimum
- The third signal: the share of requests carrying each of the 52 tags against the validation or training share
- With 52 tags, chance alone reaches PSI 0.2 in a small window, so the tag bar is about 0.39 at 200 requests
- A forty-character ticket fires 18 of the 52 tags: the lowest thresholds from the card at work

# The column you sort on is not the model

Two epochs: test P0 recall 0.806, higher than the served 0.742. The gate refuses to run before the export: `export_report.json is missing`.


Backtest on 794 tickets, fp32 as served: priority agreement 0.720, tag set agreement 0.029, 29 P1 tickets moved to P0. `GATE FAILED`: tag micro-F1 0.375, both heads down against production, both languages down.

## Retrain with a tightened bar

`--min-tag-micro-f1 0.55` is a pipeline parameter, not a code change. The run ends `FAILURE` at the register step; nothing reaches the registry.


`retrain-semantic.yaml` is on demand only: a hosted CPU runner is too slow for a weekly habit.

## On the laptop: a checkpoint, two graphs and a second service



The LoRA checkpoint exports to ONNX; the semantic engine loads the graph the gate chose, beside the triage service and its model. The same on every track.

## On the platform: a second pipeline, registry entry and endpoint (AWS)



SageMaker Pipelines registers the semantic candidate in the Model Registry; the approval deploys your semantic serverless endpoint beside triage. The EventBridge schedule ships off.

# Same loop, four platforms

| Component      | AWS                                                 | Google Cloud                         | Azure                                     | Local                      |
|----------------|-----------------------------------------------------|--------------------------------------|-------------------------------------------|----------------------------|
| Pipelines      | SageMaker Pipelines                                 | Pipelines, the Local YAML            | Azure ML pipelines, serverless            | Kubeflow local runner      |
| Registry       | package group - <code>semantic</code>               | Model Registry, aliases              | Azure ML registry, <code>stage</code> tag | MLflow                     |
| Tenant serving | Serverless Inference                                | Cloud Run, <code>NW_MODEL_URI</code> | managed online endpoint                   | the course service         |
| Drift alarm    | <code>drift_alert</code> metric filter              | log-based metric, alert policy       | log alert                                 | Evidently report, no alert |
| Retraining     | <code>pipeline-submit</code> with <code>ARGS</code> | same                                 | same                                      | same, locally              |

# Acceptance criteria

---

1. Trainable parameters under 2 percent; the data passed its contract
2. ONNX fp32 within  $1e-4$  of PyTorch; int8 smaller; four benchmark rows measured
3. The gate read three files, chose the served graph and moved `latest`
4. The pipeline registered one candidate; you approved it with a reason
5. The service serves that version from the registry, in the format the gate chose

# 40 passed

what `make session03` prints in the solution when you are done

# Questions worth asking

---

- Project 2 is slower than Project 1 by a factor you measured. For which tickets is that worth paying?
- The German tickets are five percent and the base model is English. What would you change, and what does it cost per thousand tickets?
- The gate refused int8 over three emergencies. Who should own that bar, and what would make int8 servable?
- If the customer's stack is TensorFlow Serving, what in this part transfers unchanged?

RESERVE

# Small models in the wider market, and the other tracks

---

Where small and on-device models are going, then the platform band for Google Cloud, Azure and Local.

# Vendors now ship small models for phones and the edge, and Gartner predicts wider use

---

## A prediction

Gartner predicts that by 2027 organizations will use small, task-specific models at least three times more than general-purpose LLMs, by usage volume.

## On the phone


Apple's on-device model has about 3B parameters at 2 bits per weight, open to developers through its Foundation Models framework since 2025.

## Open and at the edge

Gemma 4's edge models: Google says the smallest runs in under 1.5 GB of memory on some devices.


Big models set the frontier. Small ones can do a narrow job closer to the data, at a price per call you can plan.

# The platform band, Google Cloud track




Pipelines on the Agent Platform (formerly Vertex AI) register the semantic candidate in the Model Registry; the approval gives you a semantic Cloud Run service reading the approved version. Cloud Scheduler ships paused.

# The platform band, Azure track



Azure Machine Learning pipelines register the semantic candidate with a stage tag; the approval deploys your semantic managed online endpoint beside triage. The ML schedule ships off.

# The compose stack, Local track



The Kubeflow local runner registers into MLflow (runs, registry, prompts), artifacts land in RustFS, and the weighted nginx proxy serves the alias. Nothing leaves the laptop.